

YIP TIN PO

Hong Kong Permanent Resident tpyip@connect.ust.hk <https://yip-tin-po-alex.github.io/>

Education

Hong Kong University of Science and Technology (HKUST)

Sep 2023 – Jun 2027

Bachelor of Computer Engineering

- **GPA:** 3.2/4.0 **English:** IELTS 8.0
- **Core Modules:** Deep Learning, Data Mining, Data Structures and Algorithms, Mathematics and Statistics

Research Interests

Computer Vision, Medical AI, Explainable AI, Time Series

Research Publications

- **Bidirectional Temporal Dynamics Modeling for EEG-based Driving Fatigue Recognition.**
Under submission to IEEE Transactions On Intelligent Transportation Systems (T-ITS) [JIF : 8.4]
Yip Tin Po*, Jianming Wang*, Yutao Miao, Jiayan Zhang, Yunxu Zhao, Xiaomin Ouyang, Zhihong Li, Nevin L. Zhang
- **LogiConBench: Benchmarking Logical Consistencies of LLMs.**
ICLR 2026 (Poster)
Zheng CHEN*, Chuan Zhou*, Fengxiang Cheng, **Yip Tin Po**, Fenrong Liu, Yisen Wang, Jiajun Chai, Xiaohan Wang, Guojun Yin, Wei Lin, Haoxuan Li, Bo Li, Zhouchen Lin
- **Good Arguments Against the People Pleasers: How Reasoning Mitigates (Yet Masks) LLM Sycophancy.**
ACL 2026 (Main Conference)
Zhaoxin Feng, Zheng CHEN, Jianfei Ma, **Yip Tin Po**, Emmanuele Chersoni, Bo Li
- **PCA-guided Activation Scaling for Monotonic Bidirectional Control over LLM Sycophancy.**
Under submission to COLM 2026
Zheng CHEN, Zhaoxin Feng, **Yip Tin Po**, Jianfei Ma, Emmanuele Chersoni, Bo Li

Research Experience

Machine Learning Research Assistant

Aug 2025 – Apr 2026

Driving Fatigue Detection

- Engineered an automated facial analysis pipeline to generate standardized annotations for supervised learning
- Independently led a T-ITS publication on EEG-based driving fatigue detection, covering the entire lifecycle: dataset curation, model architecture/hyperparameter design, experimentation, and manuscript authorship (methodology, results, conclusion).
- Produced all figures, tables, and technical diagrams to ensure reproducibility.

Large Language Model Research Assistant

Aug 2025 – Feb 2026

LLM Explainability and Reasoning

- Implemented SFT pipelines across multiple open-source LLMs (Qwen, Gemma, Llama, etc.) and developed targeted steering/intervention mechanisms to constrain model behavior and improve reasoning consistency.
- Curated large-scale open datasets and designed a label generation pipeline to enable reproducible training with rigorous benchmarking.
- Enhanced manuscript quality through technical reviews and narrative restructuring to improve clarity, accessibility, and publication readiness.

Technical Skills

Languages: Python, C/C++, Java

Data Visualization: NumPy, Pandas, Matplotlib, Seaborn

Deep Learning: PyTorch, Transformers, MoE, LoRA

Tools & Systems: Git/GitHub, Linux/Bash, L^AT_EX, Cursor